

Anti-Aging Scheduling in Single-Server Queues: A Systematic and Comparative Study



Zhongdong Liu


Center for Networked Computing (CNC)
Department of Computer & Information Sciences
Temple University

Joint work with

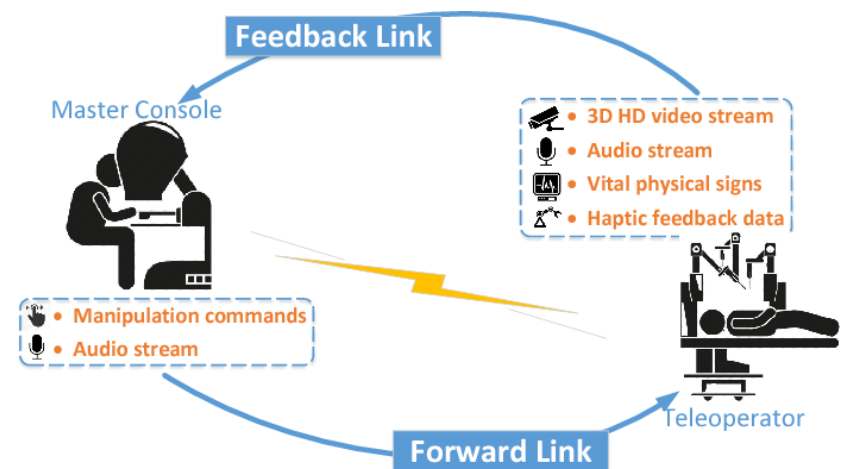
Liang Huang (Zhejiang University of Technology), **Bin Li**
(University of Rhode Island) and **Bo Ji** (Temple University)

Freshness Matters

- Real-time services are ubiquitous
 - Intelligent transportation systems & vehicular networks
 - Sensor networks (for environment/health monitoring), wireless channel feedback, news feeds, weather updates, fare aggregating, etc.




(a) Intelligent vehicular networks




(b) Sensor networks


Age of Information: Definition




- A simple abstract model
 - Source/Server/Receiver (Monitor)
 - Performance of interest:
freshness of update at the monitor



Age of Information: Definition



- A simple abstract model
 - Source/Server/Receiver (Monitor)
 - Performance of interest:
freshness of update at the monitor



- Definition


- At time t , the **Age-of-information (AoI)** $\Delta(t)$ is the “age” of the “youngest” update that was delivered to the receiver before time t
- If update i is generated at t_i and delivered at t'_i

the time elapsed since generation


the latest

$$\Delta(t) = t - \max\{t_i : t'_i \leq t\}$$


- AoI grows linearly and drops upon new update delivered



Age of Information: Definition



- A simple abstract model
 - Source/Server/Receiver (Monitor)
 - Performance of interest:
freshness of update at the monitor



- Definition


- At time t , the **Age-of-information (AoI)** $\Delta(t)$ is the “age” of the “youngest” update that was delivered to the receiver before time t
- If update i is generated at t_i and delivered at t'_i

the time elapsed since generation


the latest

$$\Delta(t) = t - \max\{t_i : t'_i \leq t\}$$

- AoI grows linearly and drops upon new update delivered




Age of Information: Definition



- A simple abstract model
 - Source/Server/Receiver (Monitor)
 - Performance of interest:

freshness of update at the monitor



- Definition

- At time t , the **Age-of-information (AoI)** $\Delta(t)$ is the “age” of the “youngest” update that was delivered to the receiver before time t

- If update i is generated at t_i and delivered at t'_i


the time elapsed since generation

the latest


AoI: defined for a sequence of updates

$$\Delta(t) = t - \max\{t_i: t'_i \leq t\}$$

- AoI grows linearly and drops upon new update delivered



AoI vs. Delay




- Low arrival rate


- Empty buffer \rightarrow low delay
- Infrequent updates \rightarrow large interarrival time & high AoI

- Large arrival rate

- Full buffer \rightarrow high delay
- Become stale while waiting \rightarrow high AoI



Aol vs. Delay




- Low arrival rate

- Large arrival rate

In M/M/1 FCFS queues: [Kaul et al., 12]

- Aol first **decreases**, then **increases** with arrival rate
- Delay **increases** with arrival rate



Aol vs Delay




Aol depends on both


- **Queueing delay** (how fast to deliver)
- **Inter-arrival time** (how often to generate)

In M/M/1 FCFS queues: [Kaul et al., 12]


- Aol first **decreases**, then **increases** with arrival rate
- Delay **increases** with arrival rate




Anti-Aging Scheduling




- Motivation & Position




Anti-Aging Scheduling




- Motivation & Position




Anti-Aging Scheduling




- Motivation & Position



Anti-Aging Scheduling




- Motivation & Position



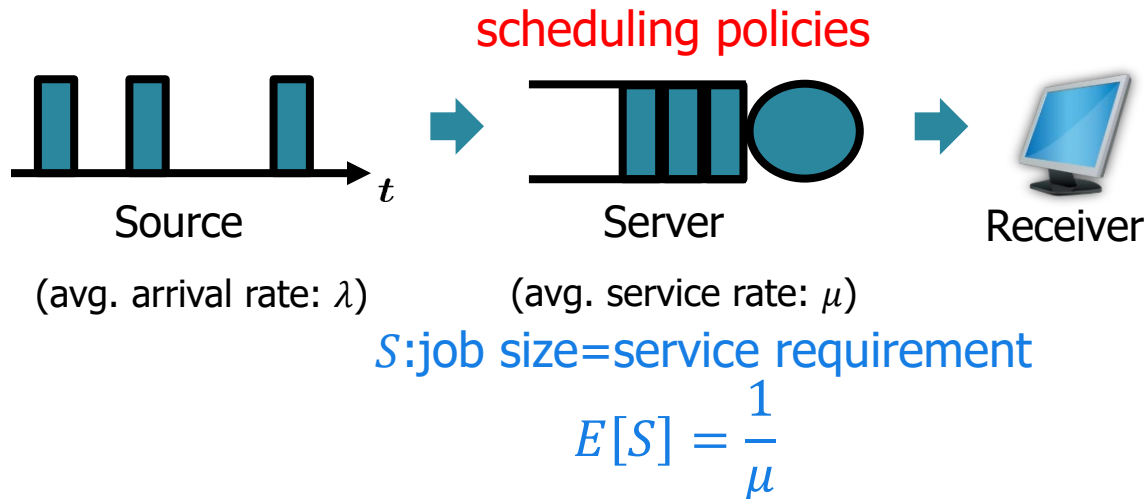
- Contributions

- Investigate the impact of scheduling policies on the AoI performance
- Summarize useful guidelines for the design of AoI-efficient policies
- Equivalence between some size-based and AoI-based policies


Anti-Aging Scheduling



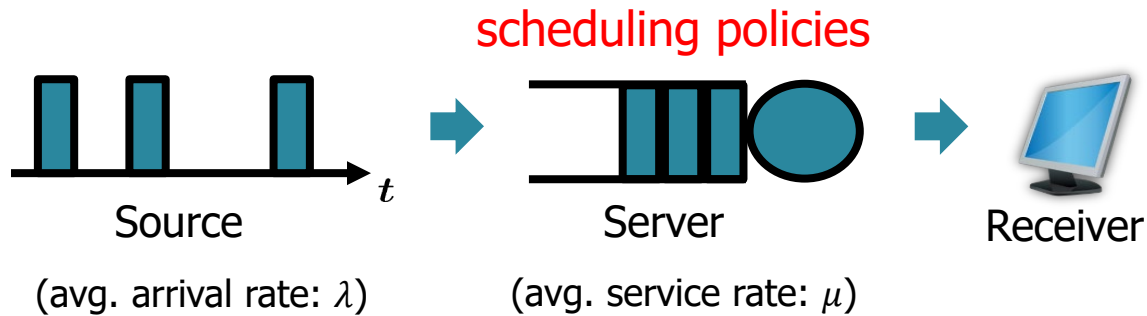
- System model: G/G/1



Anti-Aging Scheduling



- System model: G/G/1




S : job size = service requirement

$$E[S] = \frac{1}{\mu}$$


- Performance metrics

- Time Average AoI: $\Delta = \lim_{t \rightarrow \infty} \frac{1}{t} \int_0^t \Delta(\tau) d\tau$

- Average Peak AoI (PAoI): $E[A] = \frac{1}{N} \sum_{i=0}^N A_i$




Sized-based Policies



- Common scheduling policies


	Non-preemptive	Preemptive
Blind to Size	FCFS (First-Come-First-Served) LCFS (Last-Come-First-Served) RANDOM (Random-Order-Service)	PS (Processor-Sharing) LCFS_P
Uses Size	SJF (Shortest-Job-First)	SJF_P SRPT (Shortest-Remaining-Processing-Time)

Sized-based Policies




- Common scheduling policies

	Non-preemptive	Preemptive
Blind to Size	FCFS (First-Come-First-Served) LCFS (Last-Come-First-Served) RANDOM (Random-Order-Service)	PS (Processor-Sharing) LCFS_P
Uses Size	SJF (Shortest-Job-First)	SJF_P SRPT (Shortest-Remaining-Processing-Time)




Sized-based Policies




- Common scheduling policies

	Non-preemptive	Preemptive
Blind to Size	FCFS (First-Come-First-Served) LCFS (Last-Come-First-Served) RANDOM (Random-Order-Service)	PS (Processor-Sharing) LCFS_P
Uses Size	SJF (Shortest-Job-First)	SJF_P SRPT (Shortest-Remaining-Processing-Time)




Sized-based Policies




- Common scheduling policies

	Non-preemptive	Preemptive
Blind to Size	FCFS (First-Come-First-Served) LCFS (Last-Come-First-Served) RANDOM (Random-Order-Service)	PS (Processor-Sharing) LCFS_P
Uses Size	SJF (Shortest-Job-First)	SJF_P SRPT (Shortest-Remaining-Processing-Time)




Sized-based Policies




- Common scheduling policies

	Non-preemptive	Preemptive
Blind to Size	FCFS (First-Come-First-Served) LCFS (Last-Come-First-Served) RANDOM (Random-Order-Service)	PS (Processor-Sharing) LCFS_P
Uses Size	SJF (Shortest-Job-First)	SJF_P SRPT (Shortest-Remaining-Processing-Time)




Sized-based Policies




- Common scheduling policies

	Non-preemptive	Preemptive
Blind to Size	FCFS (First-Come-First-Served) LCFS (Last-Come-First-Served) RANDOM (Random-Order-Service)	PS (Processor-Sharing) LCFS_P
Uses Size	SJF (Shortest-Job-First)	SJF_P SRPT (Shortest-Remaining-Processing-Time)




Sized-based Policies



- Common scheduling policies

	Non-preemptive	Preemptive
Blind to Size	FCFS (First-Come-First-Served) LCFS (Last-Come-First-Served) RANDOM (Random-Order-Service)	PS (Processor-Sharing) LCFS_P
Uses Size	SJF (Shortest-Job-First)	SJF_P SRPT (Shortest-Remaining-Processing-Time)




Sized-based Policies




- Common scheduling policies

	Non-preemptive	Preemptive
Blind to Size	FCFS (First-Come-First-Served) LCFS (Last-Come-First-Served) RANDOM (Random-Order-Service)	PS (Processor-Sharing) LCFS_P
Uses Size	SJF (Shortest-Job-First)	SJF_P SRPT (Shortest-Remaining-Processing-Time)




Sized-based Policies




- Common scheduling policies

	Non-preemptive	Preemptive
Blind to Size	FCFS (First-Come-First-Served) LCFS (Last-Come-First-Served) RANDOM (Random-Order-Service)	PS (Processor-Sharing) LCFS_P
Uses Size	SJF (Shortest-Job-First)	SJF_P SRPT (Shortest-Remaining-Processing-Time)



Sized-based Policies


- Average AoI performance (with update size info.)




$$\text{Weibull: } \mu = 1 \text{ and } C^2 = \frac{\text{Var}(S)}{E[S]^2} = 10$$

Observation 1: Size-based policies > Non-size-based policies

Sized-based Policies




- Average AoI performance (with update size info.)




$$\text{Weibull: } \mu = 1 \text{ and } C^2 = \frac{\text{Var}(S)}{E[S]^2} = 10$$

Guideline 1: Prioritizing updates with small size

Arrival-time-based Policies




- Average Aol performance (**without** update size info.)




Weibull: $\mu = 1$ and $C^2 = 10$

Observation 2: LCFS > FCFS, RANDOM; LCFS_P > PS

Arrival-time-based Policies



- Average Aol performance (**without** update size info.)




Weibull: $\mu = 1$ and $C^2 = 10$


Guideline 2: Prioritizing recent updates

[Kaul et al., 2012; Costa et al., 2016]

Preemptive Policies




- Average AoI performance (with service preemption)



Weibull: $\mu = 1$ and $C^2 = 10$

Observation 3:

Preemptive > Non-preemption




Weibull: $\mu = 1$ and $\rho = 0.7$


Observation 4:

Preemptive policies are less **sensitive** to update size variability


Preemptive Policies



- Average AoI performance (with service preemption)



Weibull: $\mu = 1$ and $C^2 = 10$




Weibull: $\mu = 1$ and $\rho = 0.7$

Guideline 3: Allowing service preemption

[Kaul et al., 2012; Bedeway et al., 2016; Najm et al., 2018]


Current Guidelines



- Summary of guidelines

	Non-preemptive	Preemptive
Blind to Size	Guideline 1: Prioritizing updates with small size	
Uses Size	Guideline 2: Prioritizing recent updates	
	Guideline 3: Allowing service preemption	


Current Guidelines




- Summary of guidelines

	Non-preemptive	Preemptive
Blind to Size	Guideline 1: Prioritizing updates with small size	
Aol-based Policies	Guideline 2: Prioritizing recent updates	
	Guideline 3: Allowing service preemption	
Uses Size		

Aol-based Policies




- Aol-based policies
 - Aol-Drop-Earliest (**ADE**): make next Aol drop the **earliest**
 - Aol-Drop-to-Smallest (**ADS**): making the next Aol drop to the **smallest**
 - Aol-Drop-Most (**ADM**): making the next Aol drop **most**



Aol-based Policies


- Aol-based policies vs. non-Aol-based policies



Weibull: $\mu = 1$ and $C^2 = 10$

Aol-based Policies


- Aol-based policies vs. non-Aol-based policies



Weibull: $\mu = 1$ and $C^2 = 10$

Guideline 4: Aol-based policies can further improve the Aol performance


Informative Policies



- Informative policies
 - **Informative policies** only serve informative updates (can make Aol drop)
 - Almost all introduced policies have “informative” version

Informative Policies


- Informative policies
 - **Informative policies** only serve informative updates (can make Aol drop)
 - Almost all introduced policies have “informative” version
- Informative policies vs. Non-informative policies



Weibull: $\mu = 1$ and $C^2 = 10$

Informative Policies


- Informative policies
 - **Informative policies** only serve informative updates (can make Aol drop)
 - Almost all introduced policies have “informative” version
- Informative policies vs. Non-informative policies




Guideline 5: Prioritizing informative updates

[Costa et al., 2014; Pappas et al., 2015]

Informative Policies



- Informative policies
 - **Informative policies** only serve informative updates (can make Aol drop)
 - Almost all introduced policies have “informative” version
- Informative policies vs. Non-informative policies




Guideline 5: Prioritizing informative updates

[Costa et al., 2014; Pappas et al., 2015]


- Theoretical result

Proposition 1: In the G/M/1 queueing system, the Aol under LCFS_I is stochastically smaller than that under LCFS.

Preemptive, Informative, Aol-based Policies



- Simulation results



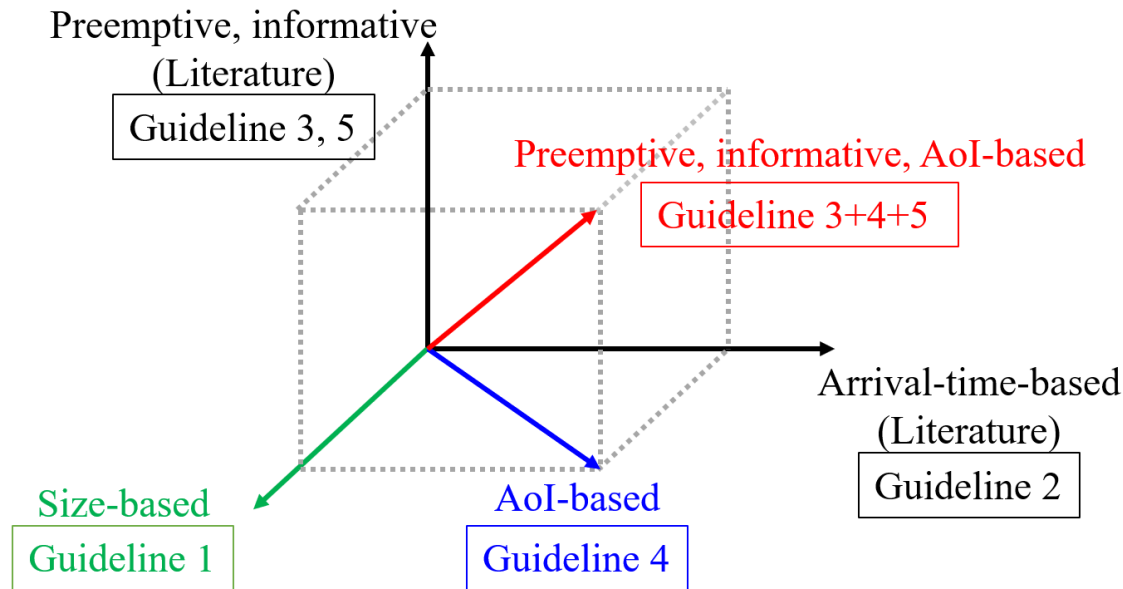

Weibull: $\mu = 1$ and $C^2 = 10$

- Sample path equivalence results

Proposition 2: SRPT_I is equivalent to ADE_PI.


Proposition 3: SJF_I is equivalent to ADE_I.

Conclusion



- Design AoI-efficient scheduling policies
 - Prioritizing small updates, allowing service preemption, prioritizing informative updates
- Equivalence between some size-based and AoI-based policies
 - $SRPT_I = ADE_{PI}$; $SJF_I = ADE_I$
- Delay-efficient \rightarrow AoI-efficient (for exogenous source)
 - High load: *delay* dominates Low load: *interarrival* dominates ⁴¹

Future Work



- Pursue more theoretical results
 - Does any informative policy always outperform its non-informative counterpart?
 - Can we derive the closed-form formulas of the average AoI/PAoI for the AoI-efficient scheduling policies (such as SRPT)?
- Apply to more complex network
 - Does our guidelines hold for multi-server queues?
 - What's the performance of our policies/guidelines in a more complex network (e.g., multi-hop networks)?



*Thank
you*

